

1. Write a parallel or distributed program that counts the number of subsets of k out of N that satisfy a given property. You have a function (**bool pred(vector <int> const& v)**) that verifies if a given subset satisfies the property. Your program shall call that function once for each subset of k elements and count the number of times it returns true.

```
#include <iostream>
#include <fstream>
#include <vector>
#include <thread>
#include <atomic>

using namespace std;

atomic <int> cnt;

inline bool check(vector <int> const &v) {
    if(v.size() == 0) {
        return false;
    }
    return v[0] % 2 == 0;
}

inline void generate(vector <int> v, int k, int n, int T) {
    if(v.size() == k) {
        for(auto it : v) {
            cerr << it << ' ';
        }
        cerr << '\n';
        if(check(v)) {
            cnt++;
        }
        return ;
    }
    int lst = 0;
    if(v.size() > 0) {
        lst = v.back();
    }
    if(T == 1) {
        for(int i = lst + 1; i <= n; ++ i) {
            v.push_back(i);
            generate(v, k, n, T);
            v.pop_back();
        }
    } else {
        thread t([&]() {
            vector <int> newv(v);
            for(int i = lst + 1; i <= n; i += 2) {
                newv.push_back(i);
                generate(newv, k, n, T / 2);
                newv.pop_back();
            }
        });
        vector <int> aux(v);
        for(int i = lst + 2; i <= n; i += 2) {
            aux.push_back(i);
            generate(aux, k, n, T - T / 2);
            aux.pop_back();
        }
    }
}
```

```

        t.join();
    }
}

int main() {
    generate(vector <int> (), 3, 5, 2);
    cout << cnt << '\n';
}

```

2. [ALREADY GIVEN] Consider the following code for enqueueing a work item to a thread pool. Find the concurrency issue and fix it. Also, add a mechanism to end the threads at shutdown.

```

class ThreadPool {
    condition_variable cv;
    mutex mtx;
    list<function<void()>> work;
    vector <thread> threads;

    void run() {
        while(true) {
            if(work.empty()) {
                unique_lock<mutex> lck(mtx);
                cv.wait(lck);
            } else {
                function<void()> wi = work.front();
                work.pop_front();
                wi();
            }
        }
    }
public:
    explicit ThreadPool(int n) {
        threads.resize(n);
        for(int i = 0; i < n; ++ i) {
            threads.emplace_back([this](){run();});
        }
    }
    void enqueue(function <void()> f) {
        unique_lock<mutex> lck(mtx);
        work.push_back(f);
        cv.notify_one();
    }
}

```

Concurrency issue

The issue here is that the ThreadPool is started with `n = 2` and on the main thread, we create another two threads that call the `enqueue()` method, both will try to access `work.front()` element, and follow with an `work.pop_front()`.

Solution

The critical resource is the list. The correct way would be to have the mutex to also lock the other branch of the `is case`

Code

```
void run() {
    while(true) {
        if(work.empty()) {
            unique_lock<mutex> lck(mtx);
            cv.wait(lck);
        } else {
            {
                unique_lock<mutex> lck(mtx);
                function<void()> wi = work.front();
                work.pop_front();
            }
            wi();
        }
    }
}
```

3. [ALREADY GIVEN] Write a parallel algorithm that computes the product of two big numbers. Or that computes the scalar product of two vectors (same thing)

```
#include <iostream>
#include <fstream>
#include <vector>
#include <thread>

using namespace std;

inline vector <int> solve(vector <int> a, vector <int> b, int T) {
    vector <int> c;
    int n = a.size();
    int m = 2 * n - 1;
    c.resize(m, 0);
    vector <thread> thr;
    for(int idx = 0; idx < T; ++ idx) {
        thr.push_back(thread([&, idx](){
            // v[i + j] += a[i] * b[j]
            // v[x] = sum(a[i] * b[x - i])
            for(int x = idx; x < m; x += T) {
                for(int i = 0; i < n; ++ i) {
                    if(x - i >= n || x - i < 0) {
                        continue;
                    }
                    c[x] += a[i] * b[x - i];
                }
            }
        }));
    }
}
```

```

    }
    for(int i = 0; i < thr.size(); ++ i) {
        thr[i].join();
    }
    return c;
}

int main() {
    ifstream fin("input.in");
    vector <int> a, b;
    int n;
    fin >> n;
    for(int i = 0; i < n; ++ i) {
        int x;
        fin >> x;
        a.push_back(x);
    }
    for(int i = 0; i < n; ++ i) {
        int x;
        fin >> x;
        b.push_back(x);
    }
    for(auto it: solve(a, b, 5)) {
        cout << it << ' ';
    }
}
}

```

3. Partial sums

Given a sequence of 'n' numbers, compute the sums of the first 'k' numbers, for each 'k' number between '1' and 'n'. Parallelize the computations, to optimise for low latency on a large number of processors. Use at most $2*n$ additions, but no more than $2*\log(n)$ additions on each computation path from inputs to an output.

Example: if the input sequence is '1 5 2 4' then the output should be '1 6 8 12'.

```

#include <iostream>
#include <fstream>
#include <vector>
#include <thread>
#include <assert.h>

using namespace std;

const int maxlg = 22;
const int maxn = 500005;

int n;
long long dp[maxlg][maxn], sum[maxn], psum[maxn];

const int T = 5;

//#define debug

```

```

inline void doIt(int idx) {
    for(int i = idx; i < n; i += T) {
        int act = 0;
        int now = i + 1;
        for(int bit = 0; (1 << bit) <= now; ++ bit)
            if(now & (1 << bit)) {
                sum[i] += dp[bit][act];
                act += (1 << bit);
            }
    }
}

int main() {

    cin >> n;
    int element;
    for(int i = 0; i < n; ++ i) {
        cin >> element;
        dp[0][i] = element;
        psum[i] = psum[i - 1] + dp[0][i];
    }

    for(int k = 1; (1 << k) < maxn; ++ k) {
        for(int i = 0; i < n; ++ i) {
            dp[k][i] = dp[k - 1][i] + dp[k - 1][i + (1 << (k - 1))];
        }
    }

    vector <thread> th;
    for(int i = 0; i < min(T, n); ++ i)
        th.push_back(thread(doIt, i));

    for(int i = 0; i < th.size(); ++ i)
        th[i].join();

    for(int i = 0; i < n; ++ i) {
        cout << sum[i] << '\n';
        assert(sum[i] == psum[i]);
    }

    return 0;
}

```

1. multiplication of two polynomials having both karatsuba and sequential implementations

```

#include <iostream>
#include <fstream>
#include <vector>
#include <ctime>
#include <thread>

using namespace std;

const int T = 3;

int n;
vector <int> a, b, sol;
vector <thread> threads;

```

```

inline void work(int idx) {
    for(int i = idx; i < n; i += T) {
        // solve for position i
        for(int x = 0; x <= i; ++ x) {
            int y = i - x;
            sol[i] += a[x] * b[y];
        }
    }
}

inline void solve() {
    cin >> n;
    for(int i = 0; i < n; ++ i) {
        int x;
        cin >> x;
        a.push_back(x);
    }
    for(int i = 0; i < n; ++ i) {
        int x;
        cin >> x;
        b.push_back(x);
    }
    sol.resize(2 * n - 1, 0);
    for(int i = 0; i < min(n, T); ++ i)
        threads.push_back(thread(work, i));
    for(int i = 0; i < threads.size(); ++ i)
        threads[i].join();
    for (auto x: sol)
        cout << x << " ";
}

int main(int argc, char* argv[]) {
    solve();
}

```

2. [ALREADY GIVEN] Write a parallel or distributed program that finds all the prime numbers up to N. Hint: serially produce all the prime numbers up to \sqrt{N} and distribute them to all threads or processes.

```
#include <iostream>
#include <fstream>
#include <vector>
#include <thread>

using namespace std;

void solve(int N) {
    vector <thread> thr;

    for (int i = 3; i <= N; i++) {
        thr.push_back(thread([&, i])) {
            bool isPrime = true;
            for (int j = 2; j <= sqrt(i); j++) {
                if i % j == 0 {
                    isPrime = false;
                    break;
                }
            }
            if (isPrime) {
                print(i);
            }
        }
    }

    for(int i = 0; i < thr.size(); ++ i) {
        thr[i].join();
    }
}

int main() {
    ifstream fin("input.in");
    int n;
    fin >> n;
    solve(n);
    return 0;
}
```